



# DataStax Hands-On Workshop

Thomas Gregg, Negib Marhoul

27. Februar 2017

# Agenda

1	Quick intro to DataStax and Team
2	Key Facts DataStax Enterprise
3	Hands-On Agenda today

DataStax

# DataStax

**400+**  
Employees

**Founded in April 2010**

Santa Clara • San Francisco • Austin •  
London • Paris • Berlin • Tokyo • Sydney

**\$190M**  
Funding



(Series E – Sept. 2014)

**500+**  
Customers



**30% +** **FORTUNE 100**

**Forbes**

2016 World's Best  
100 Cloud Companies

**Gartner**

Ranked #1 in multiple operational  
database categories

# DACH Sales & Partner Team



**Thomas Gregg**  
[thomas.gregg@datastax.com](mailto:thomas.gregg@datastax.com)  
+49 173 5473452  
Enterprise Sales Manager DACH



**Gunther Schnack**  
[gunther.schnack@datastax.com](mailto:gunther.schnack@datastax.com)  
+44 794 9448969  
Enterprise Sales Manager DACH



**Jerome Ruiz**  
[jerome.ruiz@datastax.com](mailto:jerome.ruiz@datastax.com)  
+33 63112 6090  
Channel Director Western EMEA

## DACH Technical & Customer Success Team



**Bettina Swynnerton**  
[bettina.swynnerton@datastax.com](mailto:bettina.swynnerton@datastax.com)  
Customer Engagement Manager




**Robert Stupp**  
[robert.stupp@datastax.com](mailto:robert.stupp@datastax.com)  
Solution Architect



**Patrick Callaghan**  
[pcallaghan@datastax.com](mailto:pcallaghan@datastax.com)  
SWAT



**Simon Ambridge**  
[simon.ambridge@datastax.com](mailto:simon.ambridge@datastax.com)  
Solution Engineer



**Negib Marhoul**  
[negib.marhoul@datastax.com](mailto:negib.marhoul@datastax.com)  
Solution Engineer

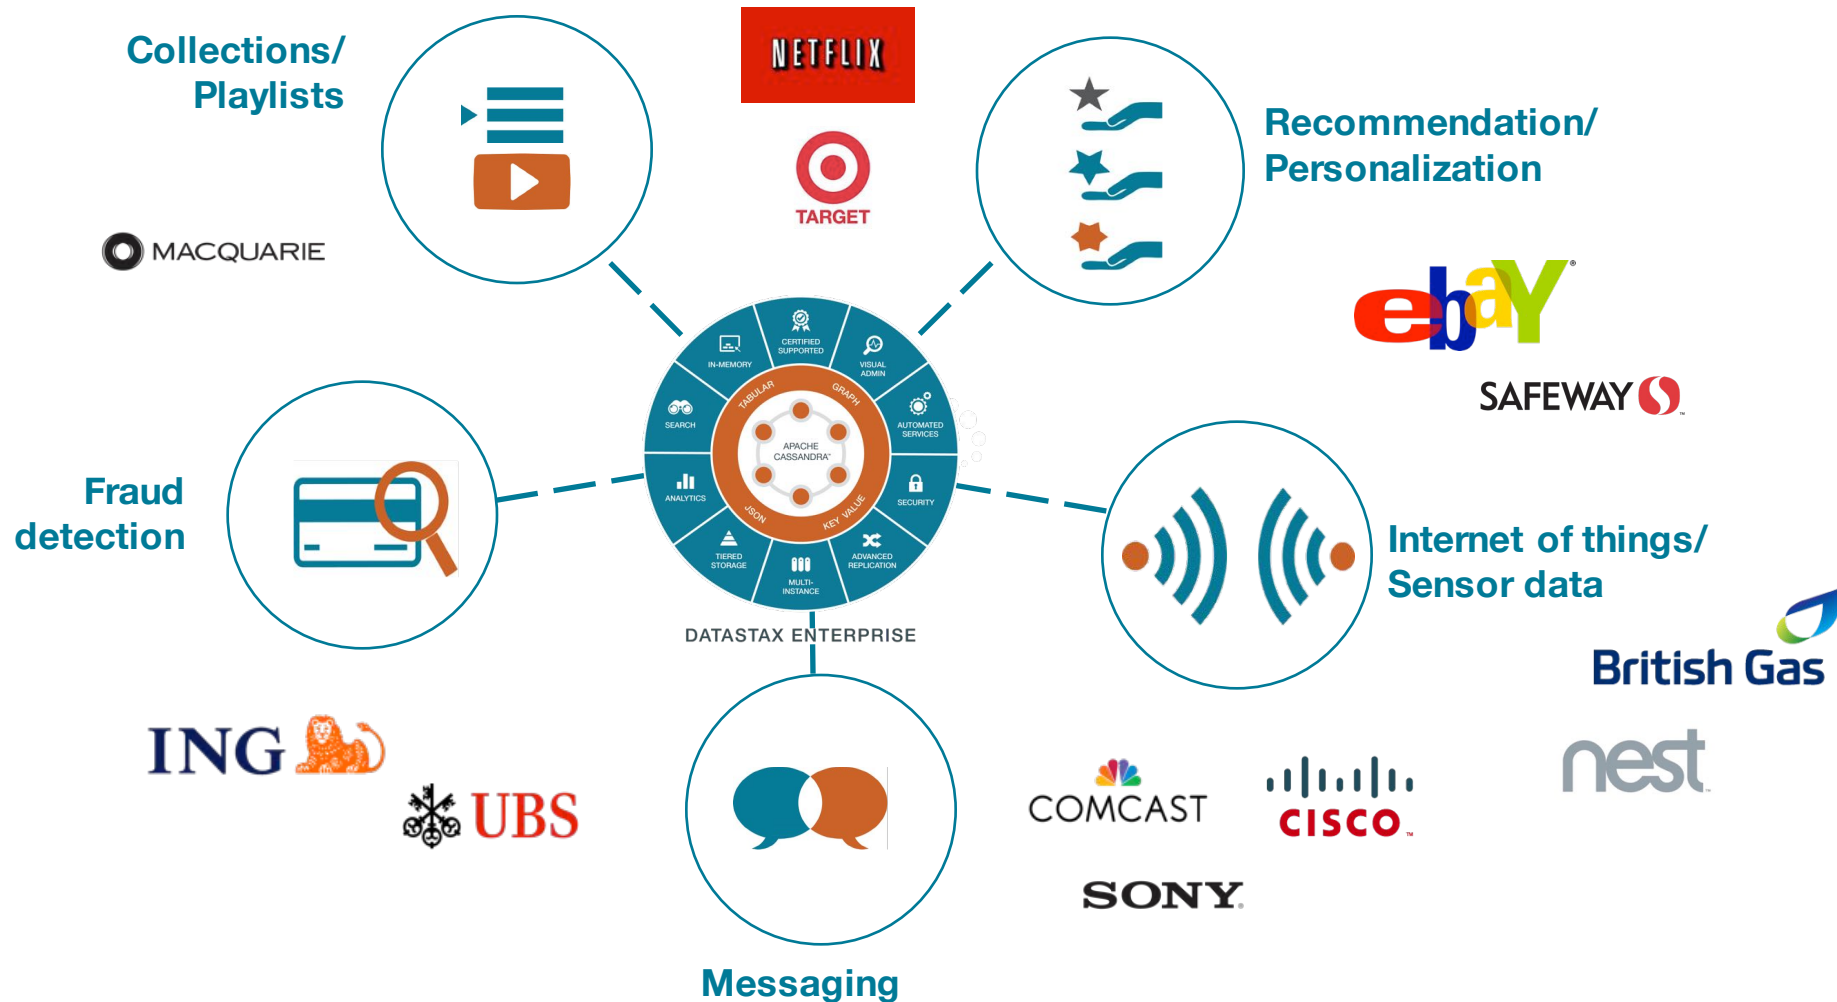


## DATASTAX ENTERPRISE

### Distributed Database Management Platform for Cloud Applications

- Always-on architecture
- Effortless Scale
- Best in class for multi-data center and cloud deployments
- Handles transactional, analytical, search, and in-memory
- workloads in a single database cluster

# Functional use cases



# DataStax Enterprise

Key Facts




# ...Driven by DataStax: The World Leader in Apache Cassandra™ Engineering

Mar 1, 2009 – Jan 8, 2017

Contributions: **Commits** ▾

Contributions to trunk, excluding merge commits




Contributors

110 Total  
22 DataStax Employees (20%)

Contributions

10,423 Total  
8,950 DataStax Employees (86%)

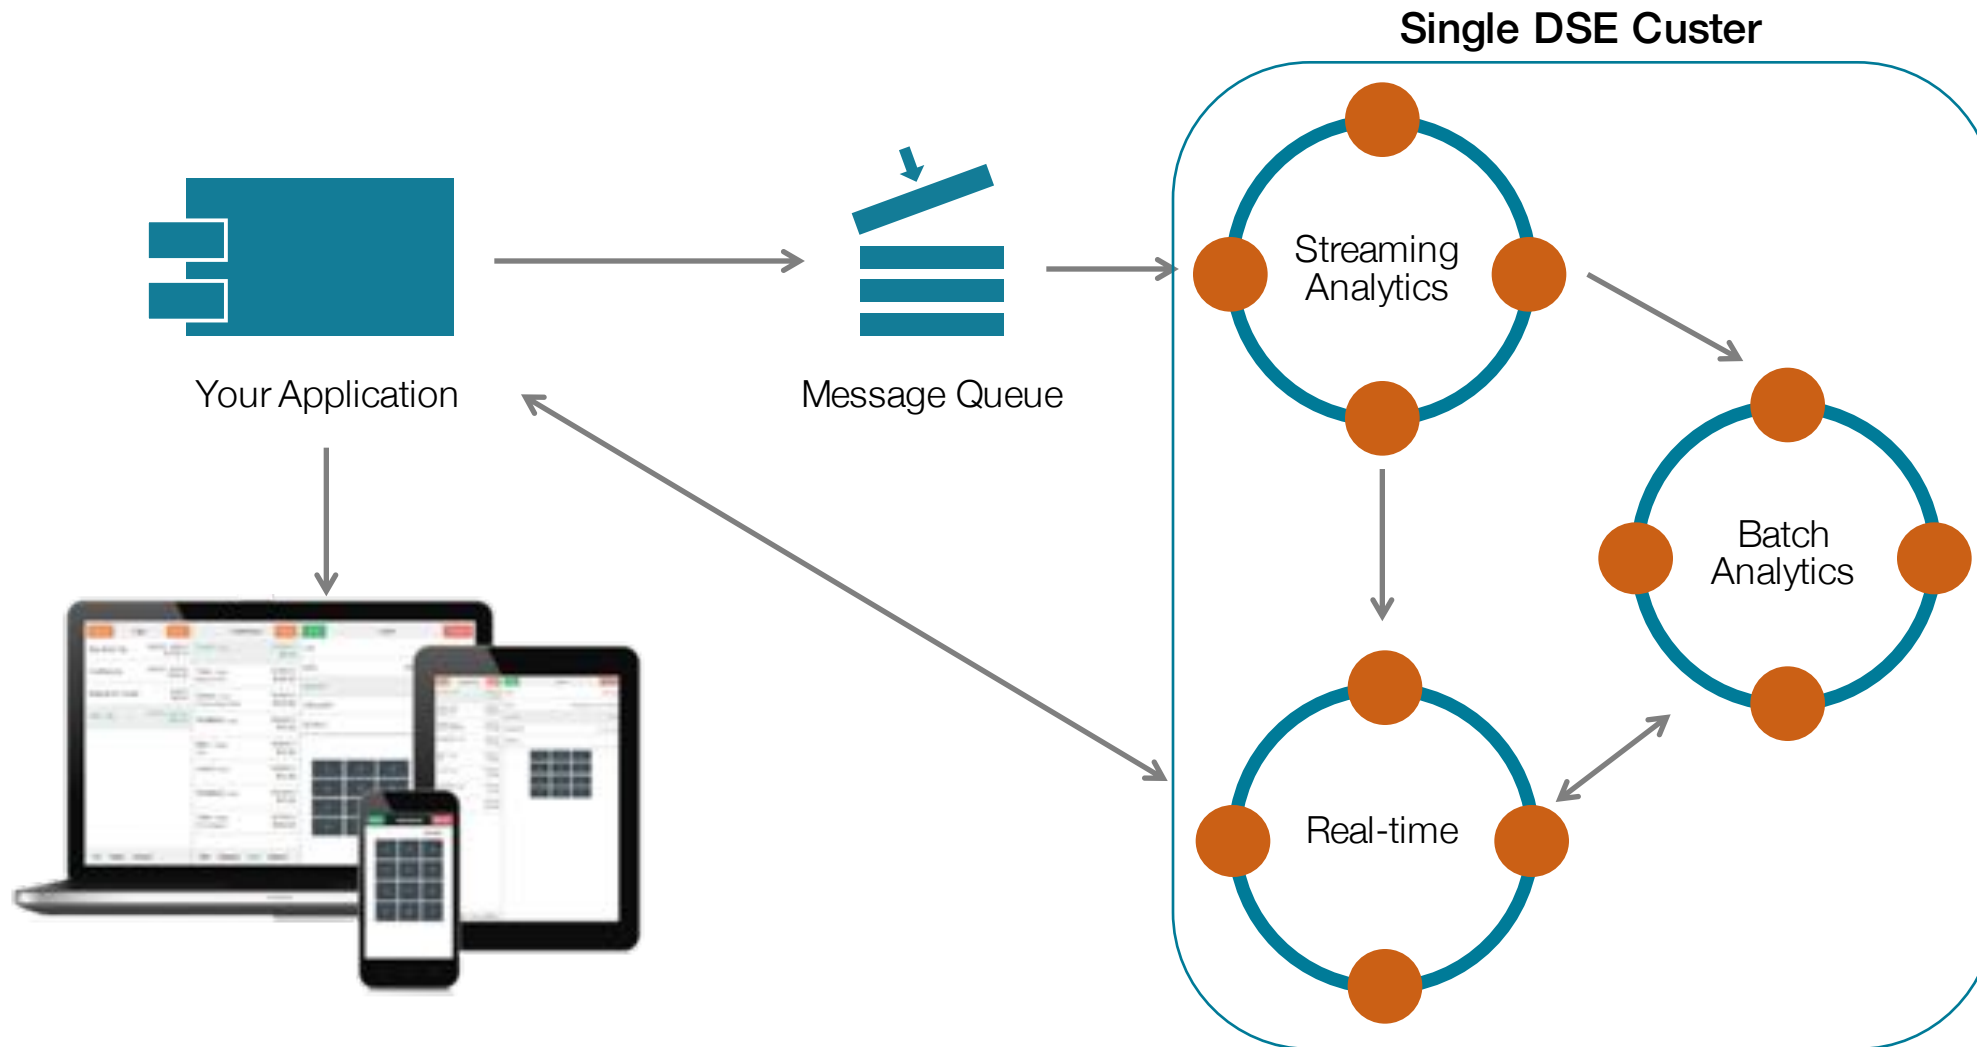
# Best distribution of Apache Cassandra™




- Continuously Available
- Master Less
- Fully Distributed
- Shared-Nothing Architecture
- Build In Replication
- Linear Scalability
- Scale out



# DSE Analytics Referenzarchitektur



# Integrated Multi-Model/Mixed Workload Platform



# DSE Search



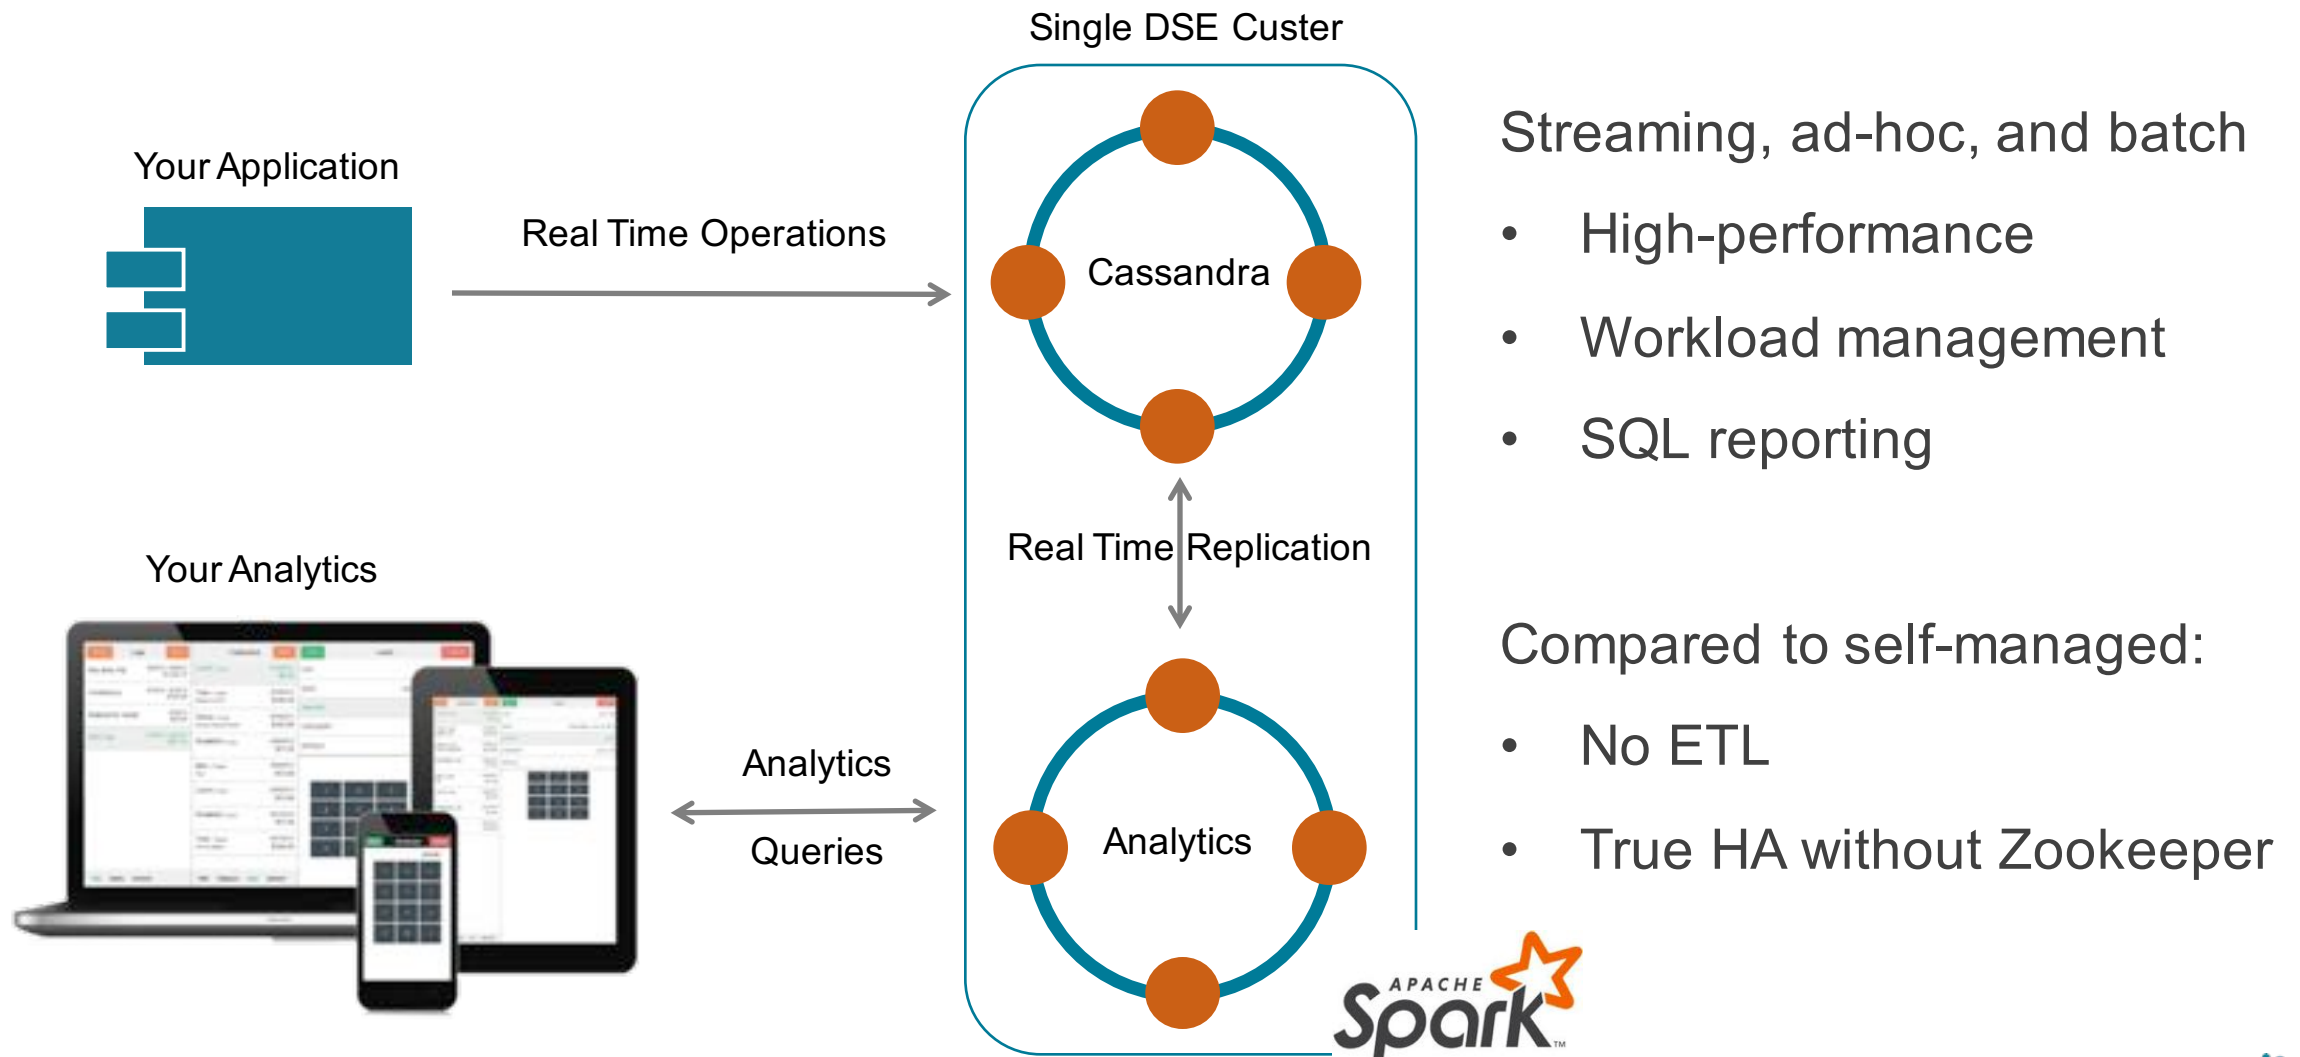
Live indexing engine with powerful search

- Automatic indexing on insert
- Higher ingestion throughput
- Distributed query optimization

Compared to self-managed:

- No separate search cluster to manage
- Probably less total hardware required
- No “Split Brain” data inconsistencies
- No ETL or synch to build and maintain
- No app level data management code


# DSE Analytics



# DSE Graph

A scalable, distributed graph database that is optimized for storing, traversing and querying complex graph data in real time

- Value data between relationships
- DSE Analytics and Search integrated
- Perfect for use cases: Customer360, Recommendations, Fraud Detection



# Agenda

1	Apache Cassandra™ Datenstrukturen und Replikationsmöglichkeiten	11:15 Uhr
2	<b>Workshop:</b> Datenmodellierung mit DataStax Enterprise und Apache Cassandra	13:00 Uhr
3	<b>Workshop:</b> Live Data Indexing mit DataStax Enterprise und Apache Solr	14:00 Uhr
4	<b>Workshop:</b> DataStax Enterprise Analytics mit Apache Spark	15:00 Uhr
5	Diskussionsrunde, Fragen & Antworten	16:00 Uhr



Vielen Dank!