



# DataStax Enterprise Architecture

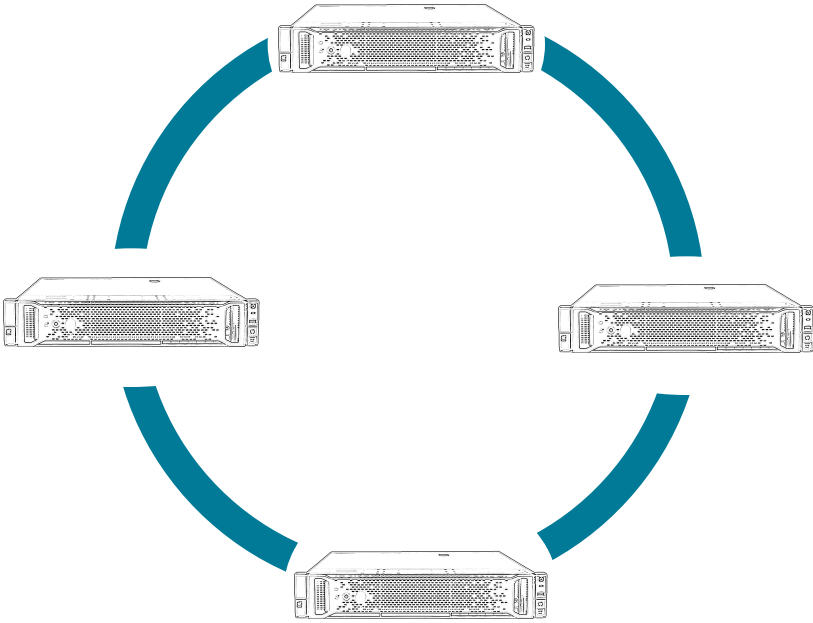
Negib Marhoul, Solution Engineer, DataStax

8. June 2017

# Agenda

|   |   |
|---|---|
| 1 | Topology and Data Structure                 |
| 2 | Request Handling                            |
| 3 | Lab1: Cassandra Access and Cassandra Stress |

# Design Goals and Objectives



- Continuously Available
- Master Less
- Fully Distributed
- Shared-Nothing Architecture
- Build In Replication
- Linear Scalability
- Scale out

Architecture

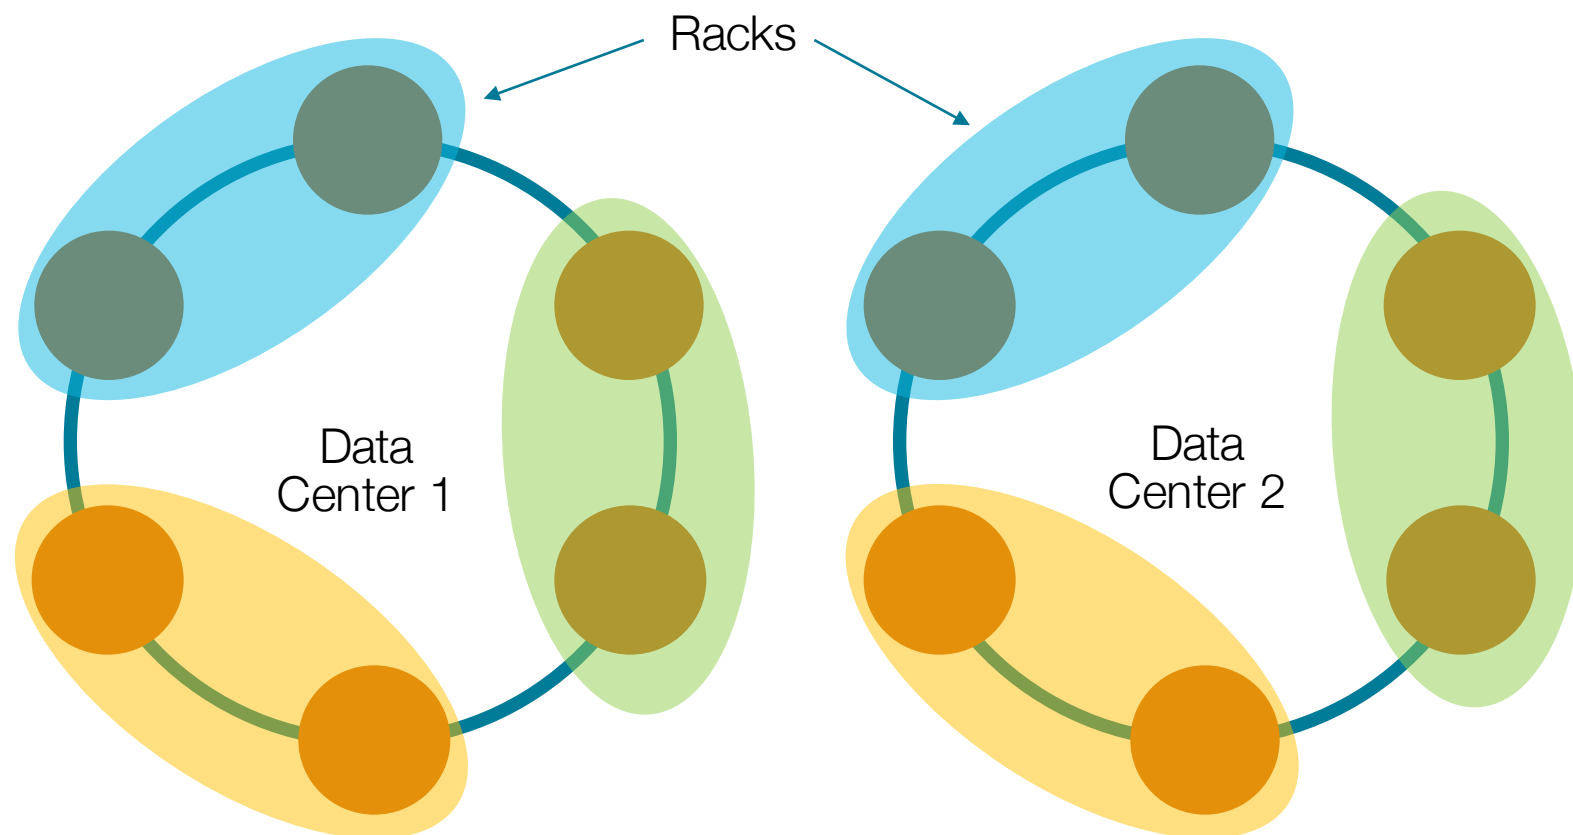
# Apache Cassandra™ Architecture

- Cluster layer
  - Amazon DynamoDB paper
  - masterless architecture
- Data-store layer
  - Google Big Table paper
  - Columns/columns family

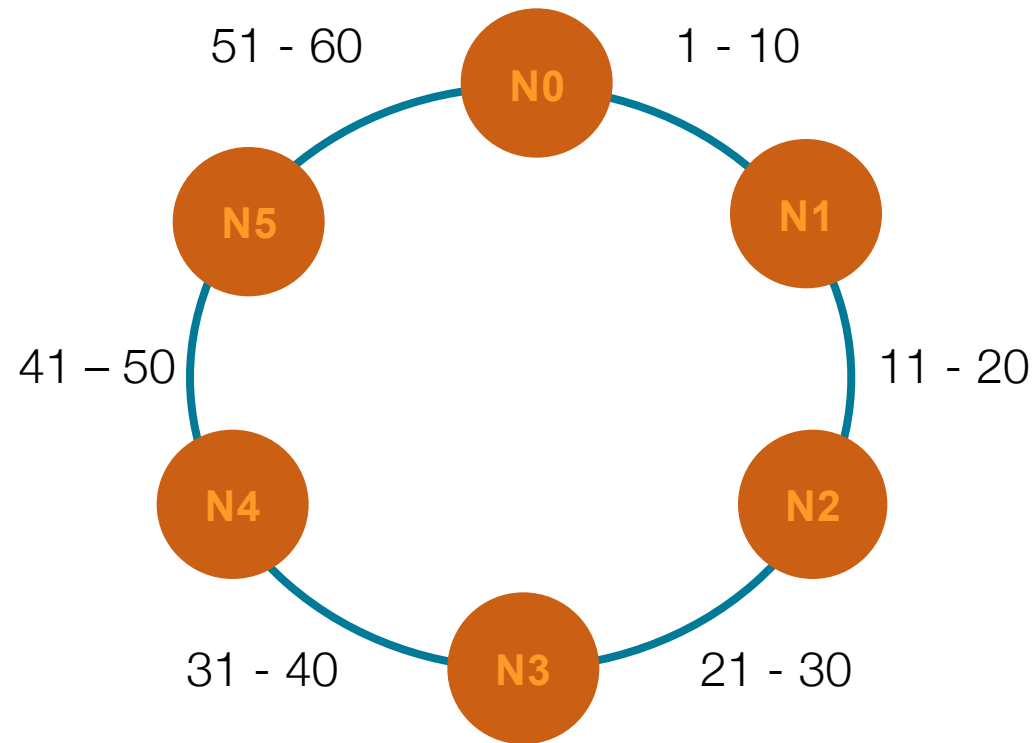


- All nodes are peers
  - Including seed nodes
  - No master
  - Discovery through gossip
- Built-in replication
  - Simplify your architecture!

## Cluster



# Token Ranges



**Token Range :  $2^{-63} - 2^{63}$**

Example with **Replication Factor 3**

N3 will own data for tokens 1 – 30

Token Range : 1-10, owned by N1, N2, N3

Token Range : 11-20, owned by N2, N3, N4

Token Range : 21-30, owned by N3, N4, N5



- Primary key
  - Partition key
  - Clustering columns
- Partitioner
  - Generates unique hash from partition key
- Replication strategy
  - Token hash determines starting point
  - Determines replica placement

# Cassandra Query Language


```
CREATE KEYSPACE retailer WITH replication =  
{'class': 'NetworkTopologyStrategy', 'DC1': '3'} ← Replication Factor  
AND durable_writes = true;
```

```
CREATE TABLE retailer.sales_by_customer (  
    orderid int,  
    salesdt date,  
    revenue double,  
    discount double,  
    comment txt, ← Partition Key  
    PRIMARY KEY(orderid);  
) WITH CLUSTERING ORDER BY (dt DESC)
```

```
SELECT * FROM sales_by_customer where custid=1 OR custid=2 AND salesdt >=20160401;
```

# Data Distribution

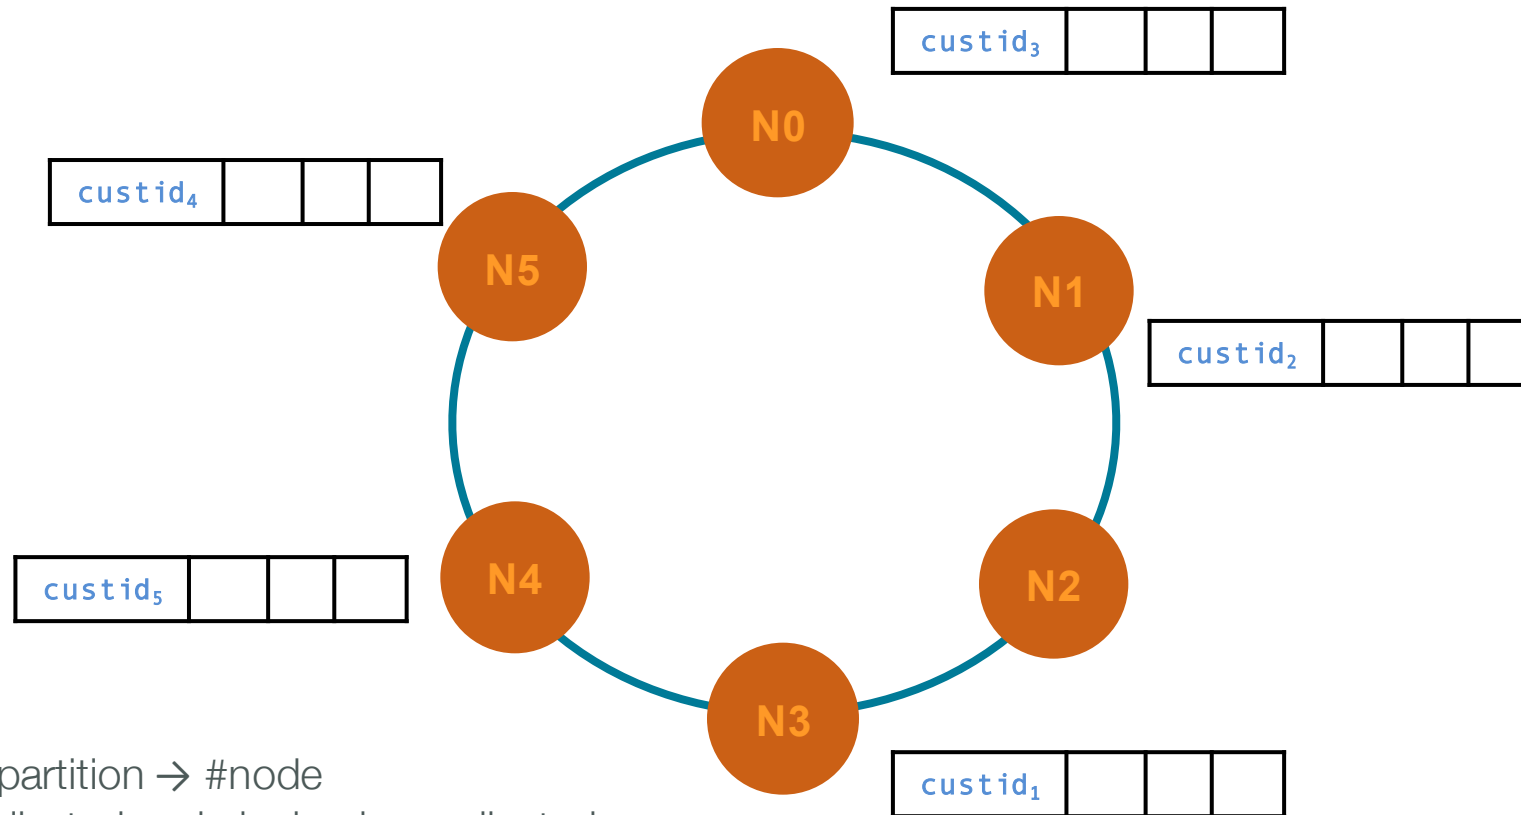
|                     |  |  |  |
|---------------------|--|--|--|
| custid <sub>1</sub> |  |  |  |
| custid <sub>2</sub> |  |  |  |
| custid <sub>3</sub> |  |  |  |
| custid <sub>4</sub> |  |  |  |
| custid <sub>5</sub> |  |  |  |



Token = hash of #partition → #node

Data is evenly distributed and clock wise replicated

# Data Distribution



Token = hash of #partition → #node

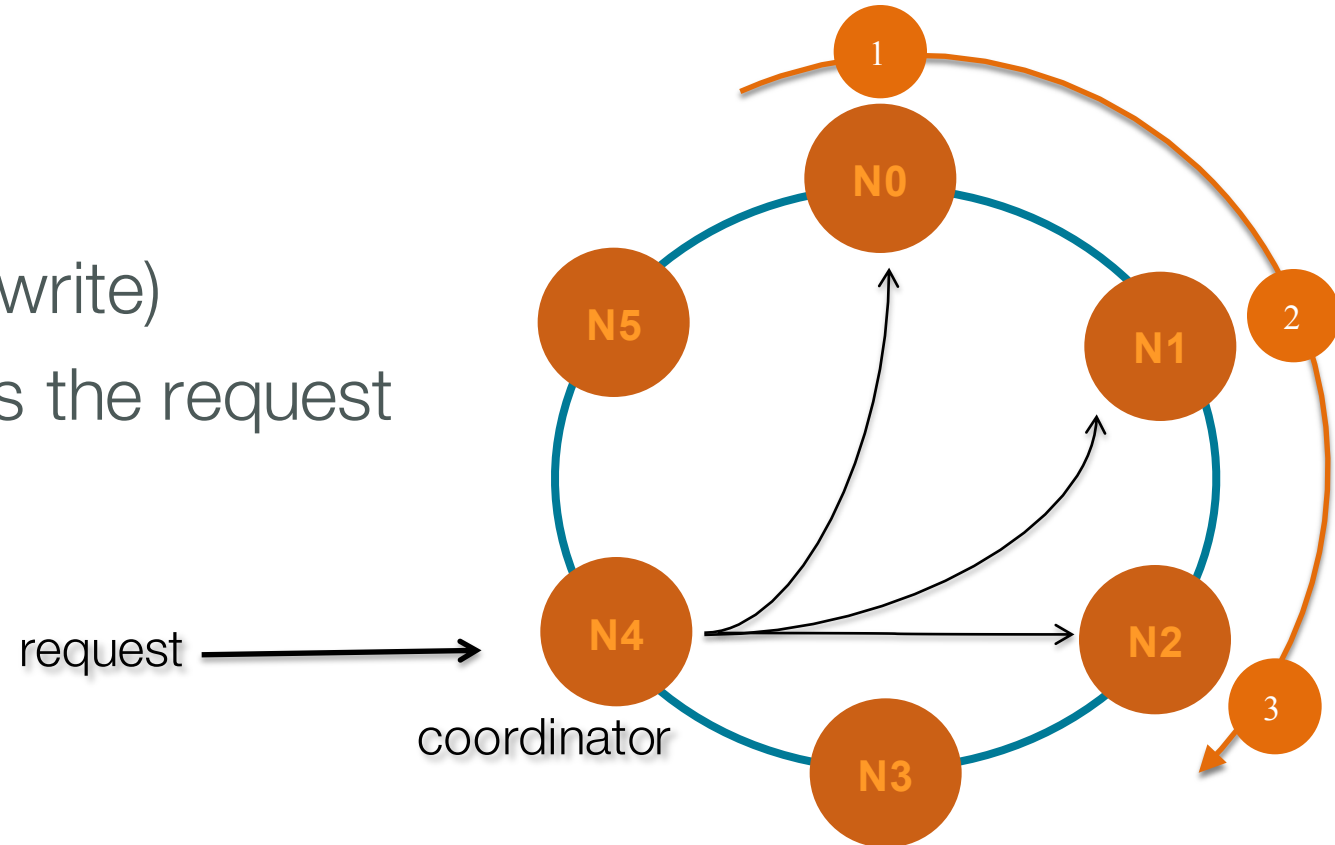
Data is evenly distributed and clock wise replicated

# Lab 1 : Accessing the cluster

Read and write request handling

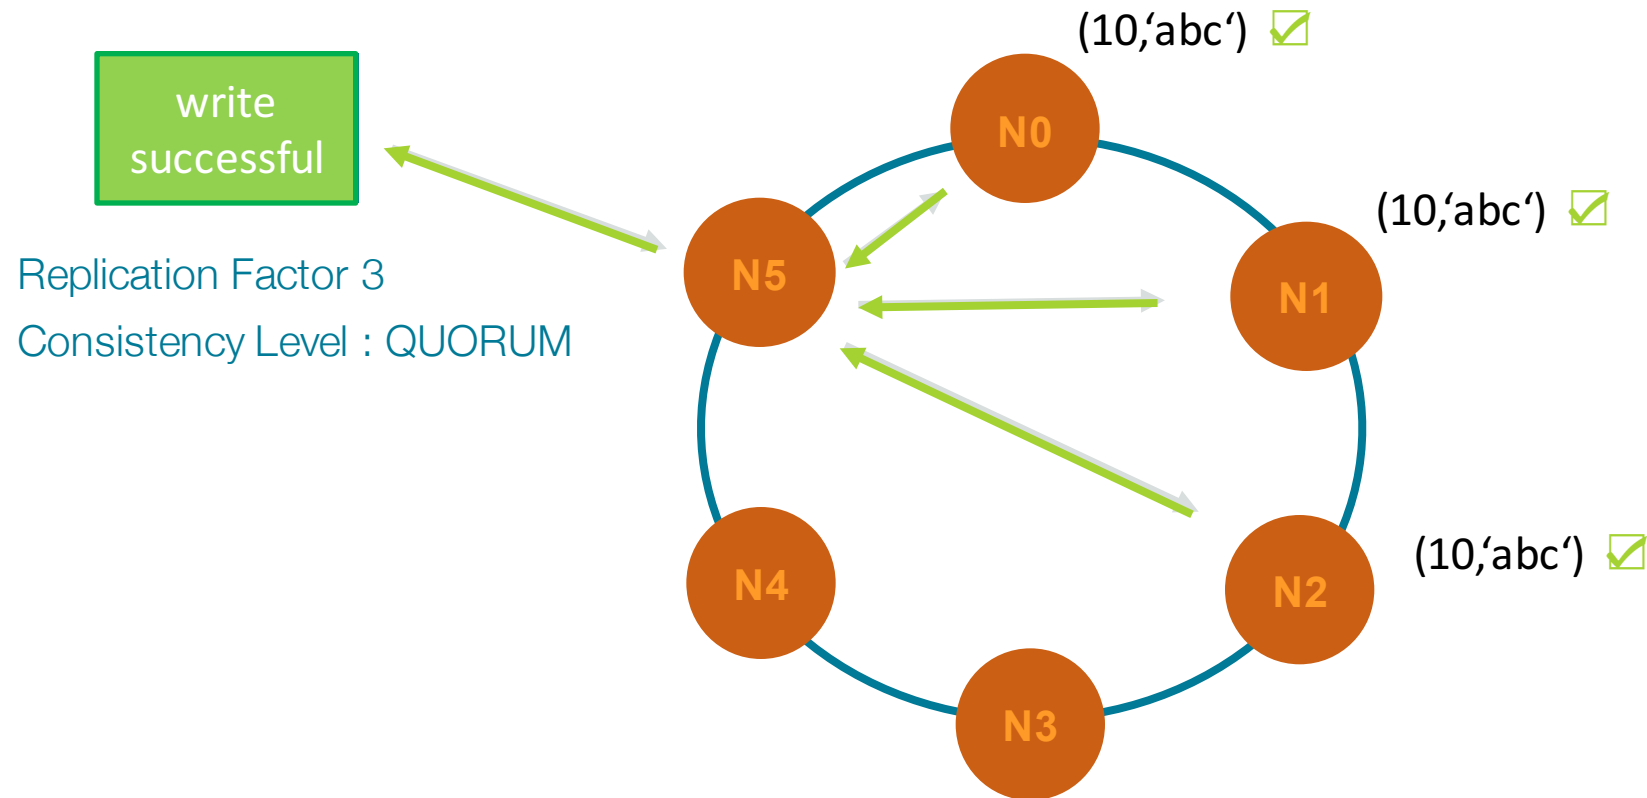
# Coordinator node

Incoming requests (read/write)  
Coordinator node handles the request



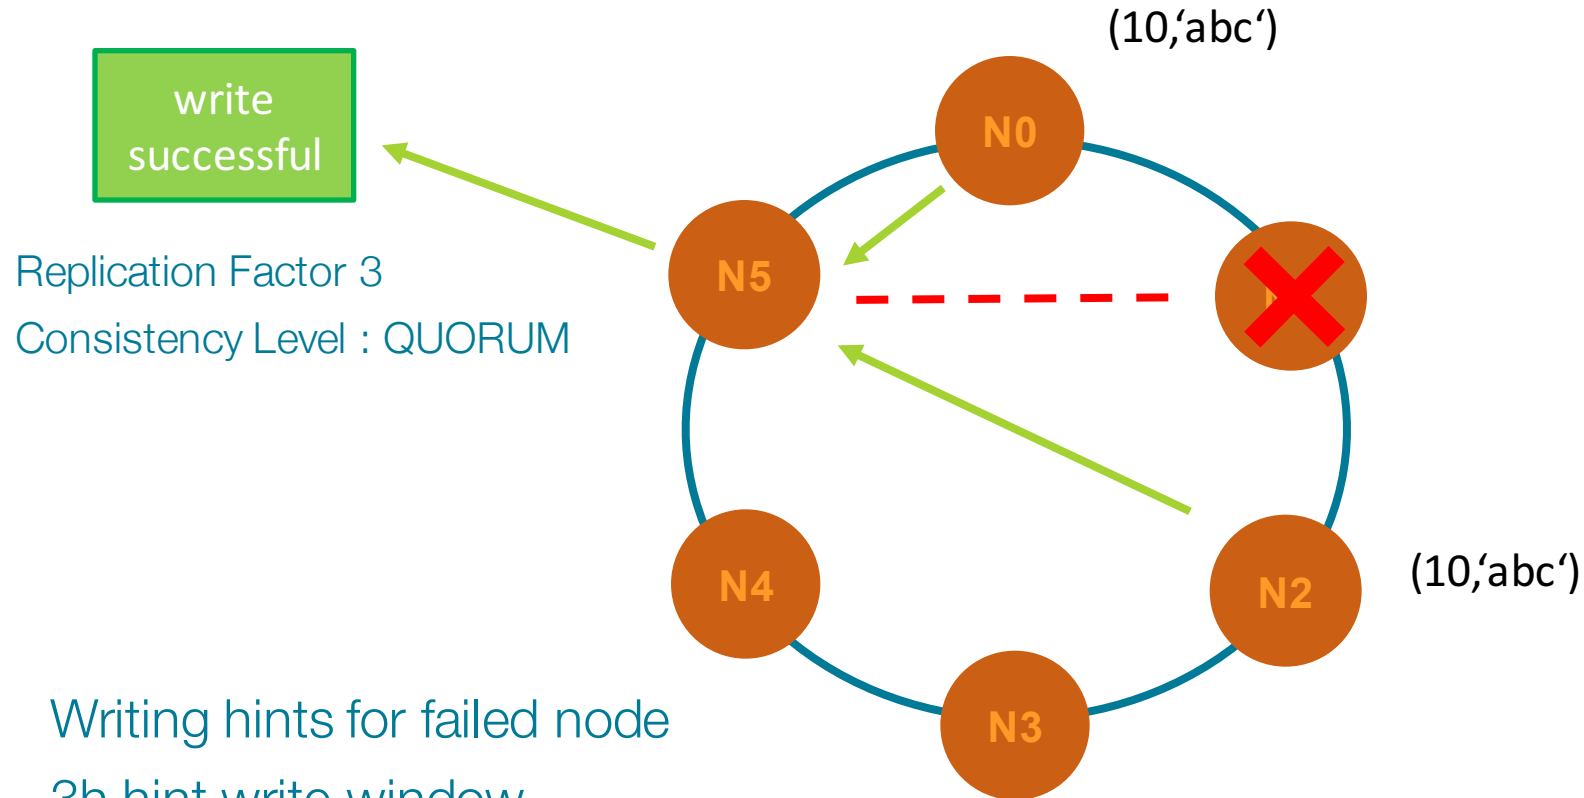
Every node can be coordinator → masterless

# Write request handling



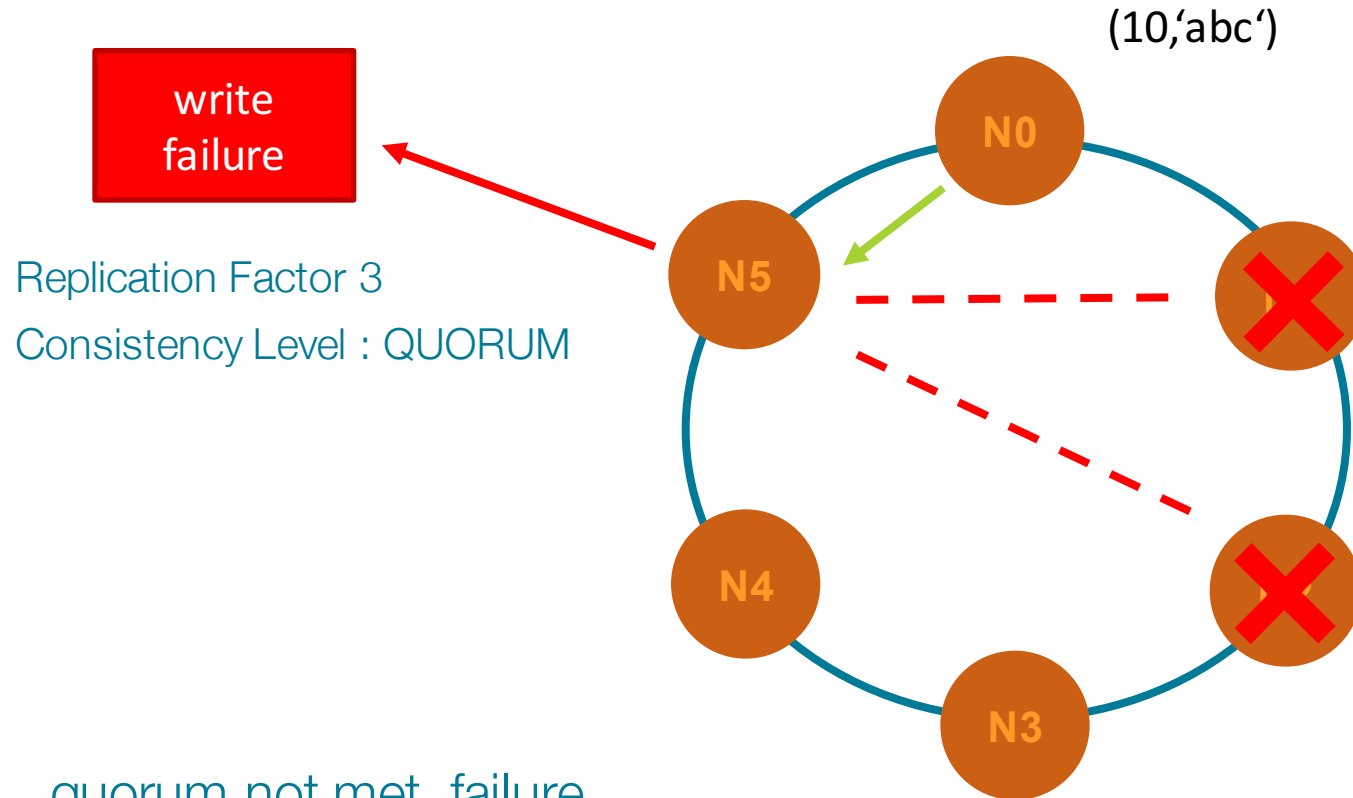


# Write request handling



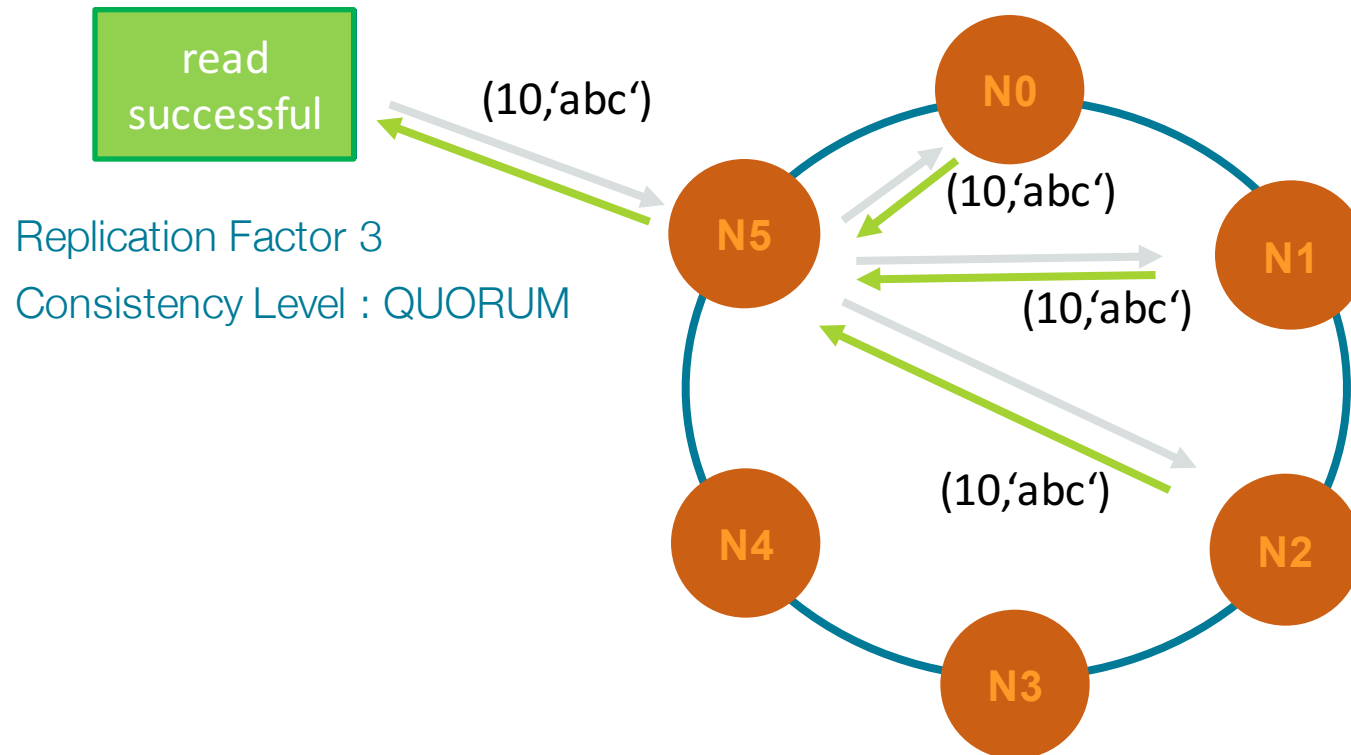
- Writing hints for failed node
- 3h hint write window
- CL: LOCAL\_ALL will fail

# Write request handling

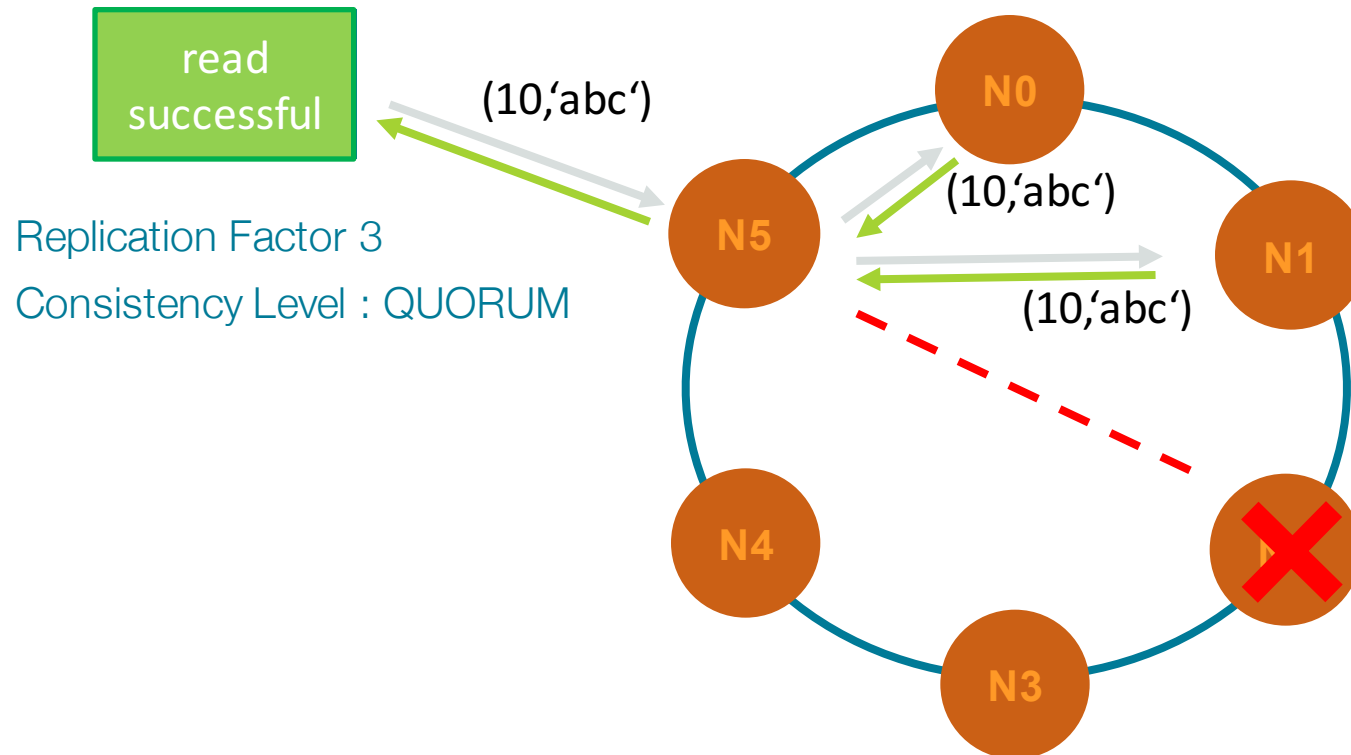


- quorum not met, failure
- CL: LOCAL\_ONE will succeed

# Read request handling

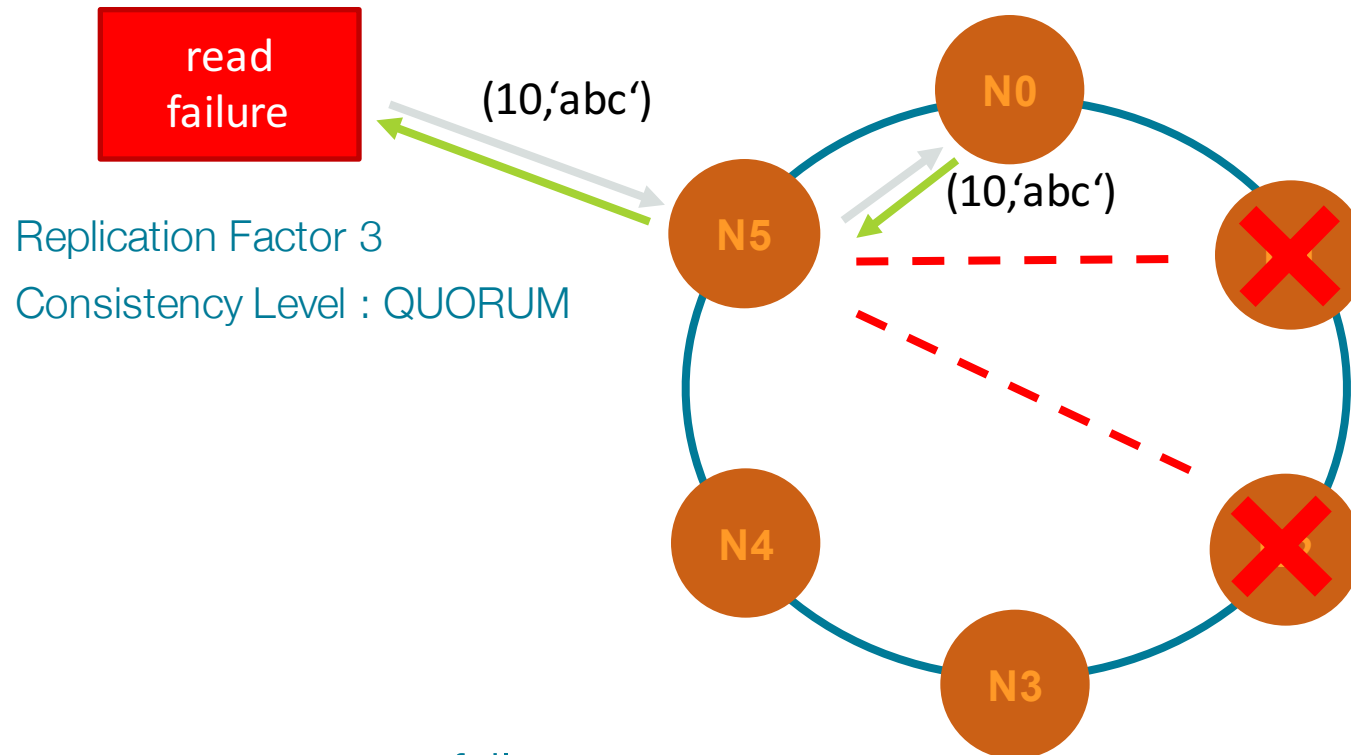


# Read request handling



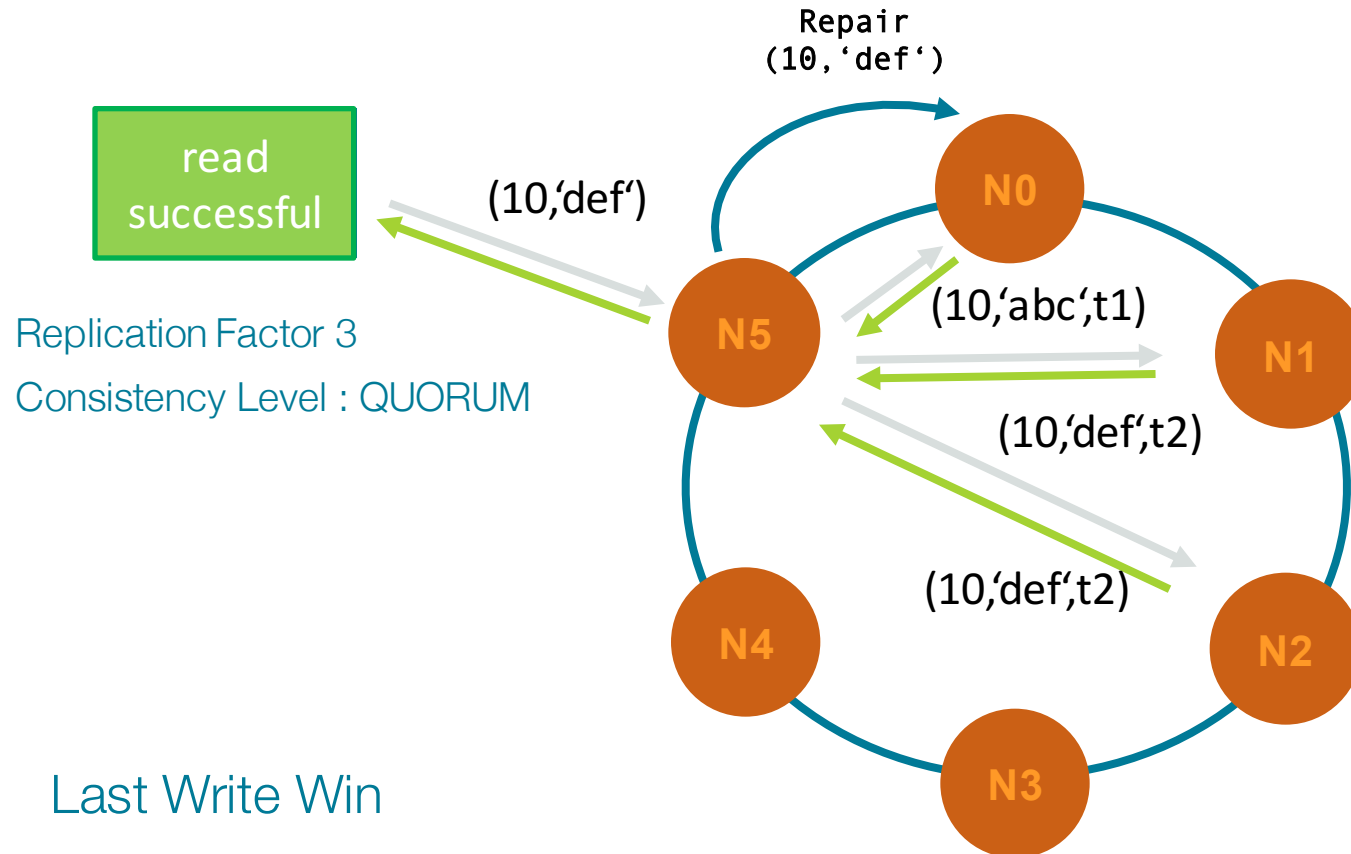
- Reading LOCAL\_QUORUM succeeds  
CL: LOCAL\_ALL will fail

# Read request handling



- quorum not met, failure
- CL: LOCAL\_ONE will succeed. [See more...](#)

# Read request handling – Read Repair



Replication Factor 3  
Consistency Level : QUORUM

- Last Write Win
- $R+W > RF$  = immediate consistency
- Background vs. foreground Read Repair. [See more...](#)

- Background vs. Foreground Read Repair
  - Compare digests
    - If any mismatch
  - re-request to same nodes (full data set)
    - compare full data sets, send update
    - block until out-of-date replicas respond
  - Return merged data set to the client
- Consistency Level
  - one, quorum, all
  - local vs. cluster wide

# Driver Code

```
Cluster cluster = Cluster.builder()
    .addContactPoint("127.0.0.1")
    .withLoadBalancingPolicy(new TokenAwarePolicy(DCAwareRoundRobinPolicy.builder()
        .withLocalDc("myLocalDC")
        .build()))
    .build();
```

```
PreparedStatement prepared = session.prepare
( "insert into sales_by_customer(custid, salesdt) values (?, ?)");
```

```
BoundStatement bound = prepared.bind("1", "20170102");
```

```
session.execute(bound); // Throws UnavailableException If consistency doesn't met, downgrade is
                           possible with corresponding RetryPolicy. Read More...
```



# Take away

- Data distribution (hash, tokens)
- Data replication (RF)
- All nodes are peer nodes , master less
- Background Read Repairs
- RetryPolicy in driver

# Lab 2 : Hands-on DSE CQL

Vielen Dank!

# Eventuell Bootstrap, ReBalance, Num Tokens VNodes

- All nodes are peers
  - Including seed nodes
  - No master
  - Discovery through gossip
- Built-in replication
  - Simplify your architecture!