



DataStax Hands-On Modelling

Negib Marhoul, Solution Engineer, DataStax

5. October 2017

Agenda

1	Storing data in Cassandra
2	Data Modelling, CQL basics
3	Hands-On Primary Keys

Storing data in Cassandra

Write request

custid : 4
name : Vincent

user_id : 4
name : Vincent

MEMTABLE

FLUSH

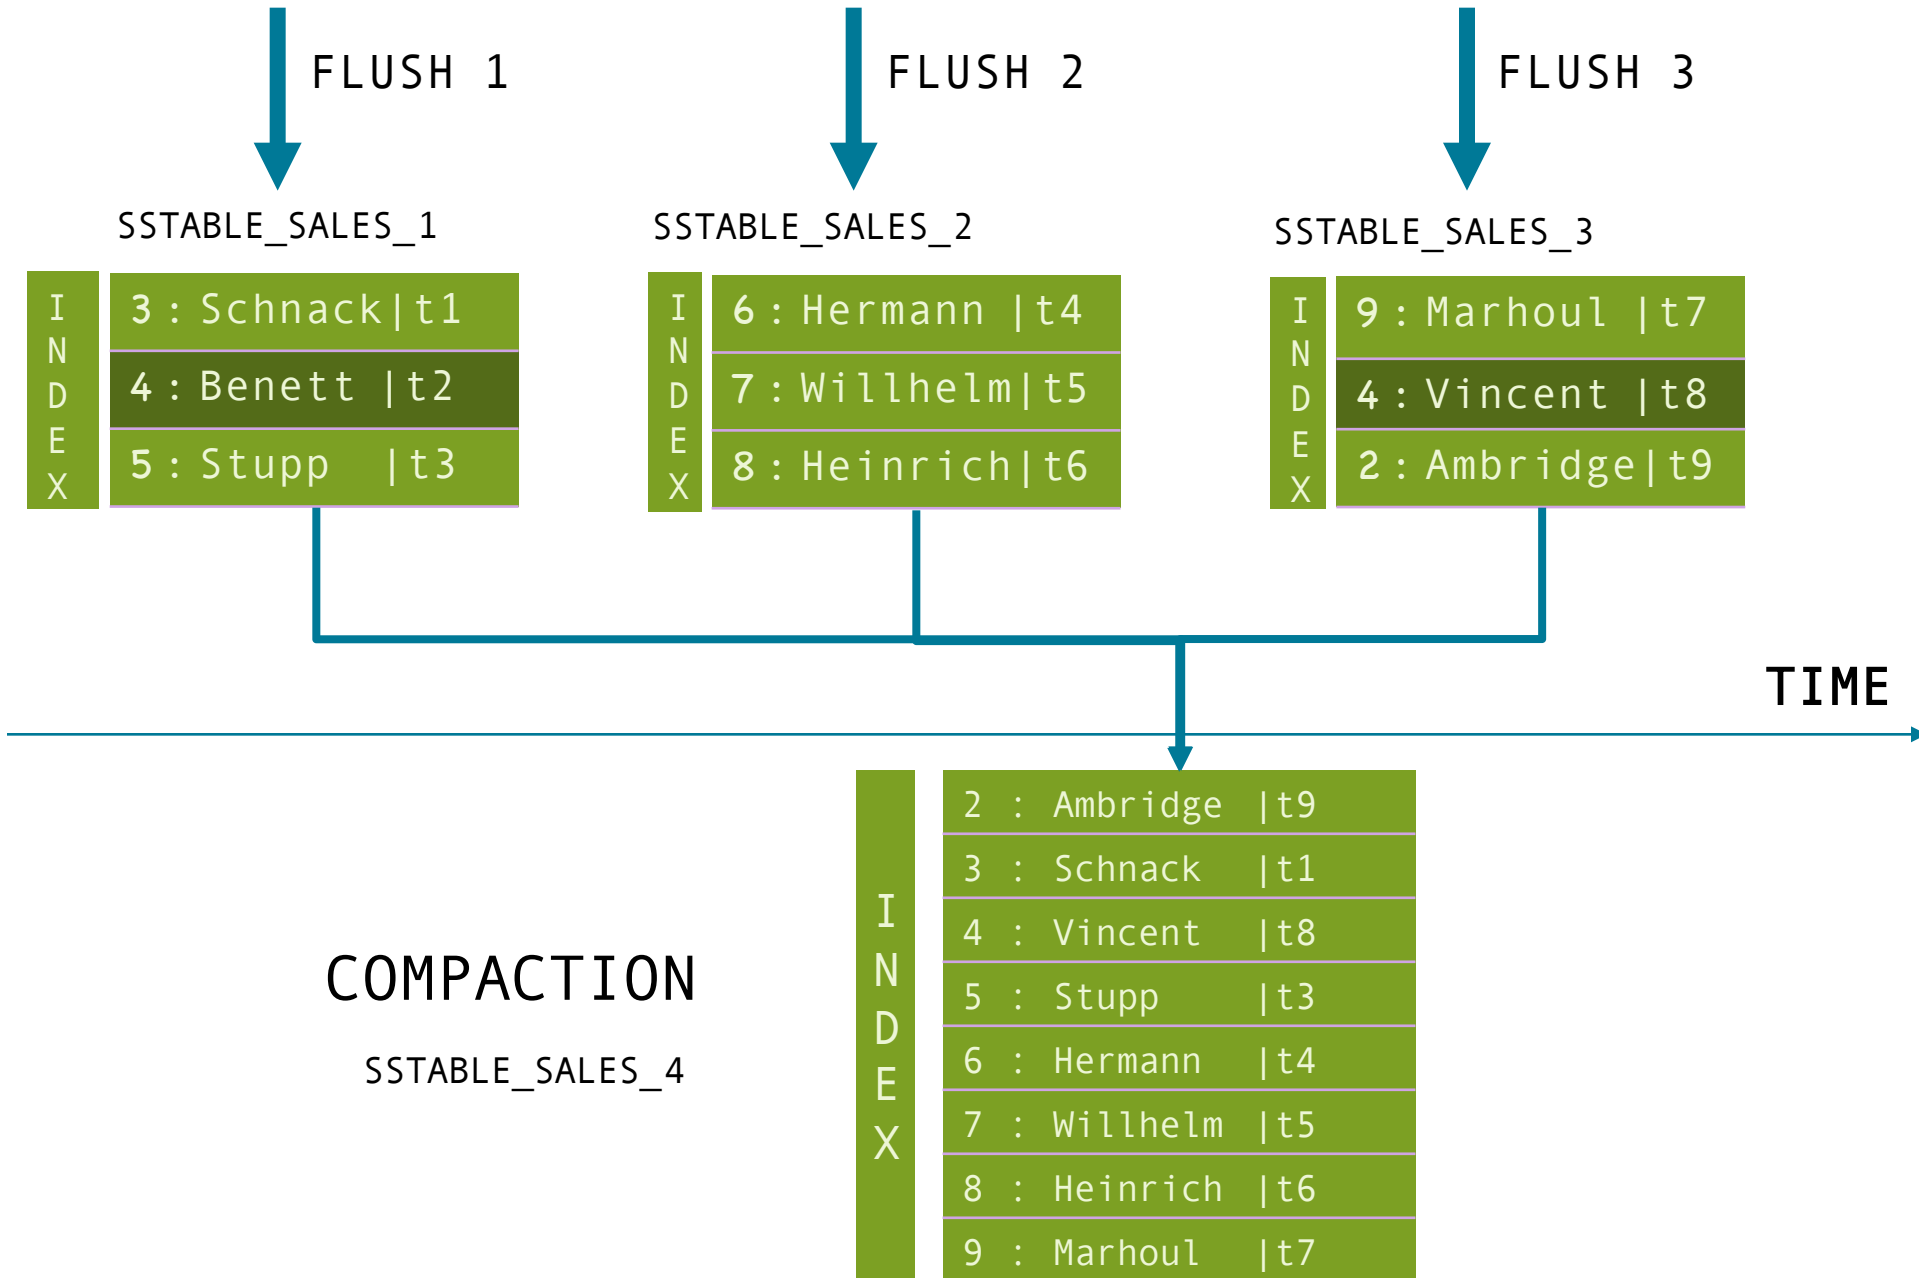
SSTABLE

I		3 : Schnack
N		4 : Vincent
D		6 : Stupp
E		8 : Ambridge
X		

IMMUTABLE
ORDERD

COMMIT LOG

user_id : 4
name : Vincent



Data Modeling Objectives

Data modeling objectives

1. Get your data **out of Cassandra**
2. Reduce query **latency**, make your queries **faster**
3. Avoid **disaster** in production

Data modeling methodology

Design by query

- first, know your **functional queries**
- then design the **table(s)** for direct access
- just **denormalize** if necessary
- Spread data evenly around the cluster
- Minimize the number of partitions read
- Make sure to take your read/update ratio into account when designing your schema

Output of design phase = **schema.cql**

Then start coding

Know your functional queries

Query:

```
find users by id group by region and orderd by join date
```

- Grouping by an attribute
- Ordering by an attribute
- Filtering based on some set of conditions
- Enforcing uniqueness in the result set

The partition key

Role

Partition key

- main entry point for query (INSERT/SELECT ...)
- help distribute/locate data on the cluster

No partition key = full cluster scan

How to choose correct partition key ?

Good partition column

- choose **functional** identifier
- high cardinality (lots of **distinct values**)

Query:

Find sales by session?


```
CREATE TABLE sales_by_customer (  
    custid int,  
    salesdt date,  
    revenue double,  
    discount double,  
    comment txt,  
    PRIMARY KEY(custid));
```

partition key (#partition)




Example of good partition key

```
CREATE TABLE sales_by_order(  
  custid int,  
  ... ,  
  PRIMARY KEY(custid))
```



Example of bad partition key

```
CREATE TABLE sales_by_cat(  
  category text,  
  ...,  
  PRIMARY KEY(category))
```



No Clustering Columns

PRIMARY KEY (custid)

// no duplicate primary keys

custid	salesdt	revenue	discount	comment
1	20160101	300	10	PlayStation4
2	20160102	699	25	iPhone 7
2	20160103	750	10	iPhone 7

```
SELECT * FROM sales_by_customer WHERE custid = '2';
```

Composite partition key

Multiple columns for partition key

- always **known in advance** (INSERT/SELECT ...)
- are **hashed together** to the same token value

```
CREATE TABLE sales_by_customer (  
    custid  int,  
    salesdt date,  
    revenue double,  
    discount double,  
    comment txt,  
    PRIMARY KEY((custid, salesdt)) );
```

```
SELECT * FROM sales_by_customer WHERE custid = ... AND salesdt = ...
```

```
SELECT * FROM sales_by_customer WHERE custid = ... AND salesdt IN (xxx, yyy ...)
```


Compound Partition Key

PRIMARY KEY ((custid, salesdt)) // hash(custid, salesdt) → token

custid	salesdt	revenue	discount	comment
1	20160101	300	10	PlayStation4
2	20160102	699	25	IPhone 7
2	20160103	750	10	IPhone 7

```
SELECT * FROM sales_by_customer WHERE custid = 2 AND salesdt = '1/2/2016';
```

The clustering column(s)

Role

Clustering column(s)

- simulate 1 – N relationship
- and sort data (logically & on disk)

Clustered table (1 – N)

```
CREATE TABLE sales_by_customer (  
  custid  int,  
  salesdt date,  
  revenue int,  
  discount int,  
  comment text,  
  PRIMARY KEY(custid, salesdt));
```

partition key

clustering column
(sorted)

unicity

Recommended syntax

```
PRIMARY KEY((custid), salesdt);
```

Clustering Columns Create Wide Rows

PRIMARY KEY ((custid),salesdt)

// default sort and range queries

custid	salesdt		
1	20160102		
	719, 10, new customer iPhone7		
2	20160101	20160102	20160103
	719, 10, new customer iPhone7	45, 50, marketing campaign xyz	187, 10, sonyplaystation 3




Merged, Sorted and Stored Sequentially

```
SELECT * FROM sales_by_customer WHERE custid = 1 and salesdt > '1/1/2016';
```

What's Stored With Each Column?

custid	salesdt		
1	20160102		
	719, 10, new customer iPhone7		
2	20160101	20160102	20160103
	719, 10, new customer iPhone7	45, 50, marketing campaign xyz	187, 10, sonyplaystation 3



column name : "comment"
column value : "new customer iPhone7"
timestamp : 1353890782373000
TTL : 3600

Columns relationship and ordering

```
CREATE TABLE sales_by_customer (  
  custid    int,  
  salesdt   date,  
  revenue   int,  
  discount  int,  
  comment   text,  
  PRIMARY KEY((custid), salesdt))  
WITH CLUSTERING ORDER BY (dt ASC)
```

Diagram illustrating column relationships and ordering:

- `custid(1) <-----> (N) salesdt` (indicated by a bracket from the `custid` and `salesdt` columns in the table definition)
- `salesdt (1) <-----> (1) (revenue, discount, comment)` (indicated by a bracket from the `revenue`, `discount`, and `comment` columns in the table definition)

custid	salesdt		
1	20160102		
	719, 10, new customer iPhone7		
2	20160101	20160102	20160103
	719, 10, new customer iPhone7	45, 50, marketing campaign xyz	187, 10, sonyplaystation 3

```
SELECT * FROM sales_by_customer WHERE custid = 1 and salesdt > '1/1/2016';
```

Multiple clustering columns

```
CREATE TABLE sales_by_cust (  
  custid    int,  
  salesch   text,  
  salesdt   date,  
  revenue   int,  
  discount  int,  
  comment   text,  
  PRIMARY KEY((custid), salesch, salesdt))  
  WITH CLUSTERING ORDER BY (salesch ASC, salesdt DESC)
```

→ salesch(1) <-----> (N) date

custid		salesch, salesdt	
1	online		
	20160102		
	719, 10, new customer iPhone7		
2	online	online	store
	20160101	20160102	20160103
	719, 10, new customer iPhone7	45, 50, marketing campaign xyz	187, 10, sonyplaystation 3

```
SELECT * FROM sales_by_customer  
WHERE custid = 1 AND  
salesch = "online" AND salesdt >= 1/2/16 AND salesdt <= 1/3/16;
```


Primary key summary

PRIMARY KEY((custid), salesch, salesdt)



Unicity of (custid, salesch, salesdt)

Primary key summary

PRIMARY KEY((custid), salesch, salesdt)



Used to locate **node** in the cluster

Used to locate **partition** in the node

Primary key summary

PRIMARY KEY((custid), salesch, salesdt)



Used to lookup rows in a partition

Used for data sorting and range queries

Lab 3 : Hands-on Primary Keys

Modelling Challenge (almost no challenge)

Query:

- Find all user details by user id
- Give me one user transaction by user id and date
- Give me all user transaction by user id between two dates
- **User details:**
 - user_id, firstname, lastname, ...
- **User transactions:**
 - user_id, date, item, quantity, price, ...

Other critical details

Huge partitions

```
PRIMARY KEY((sensor_id), dt))
```

Data for the same sensor stay in the same partition on disk

Huge partitions

PRIMARY KEY((sensor_id), dt))

Data for the same sensor stay in the same partition on disk

If insert rate = 100/sec, how big is my partition after 1 year ?

→ $100 \times 3600 \times 24 \times 365 = 3\,153\,600\,000$ cells on disks

Huge partitions

PRIMARY KEY((sensor_id), dt))

Theoretical limit of # cells for a partition = 2×10^9

Practical limit for a partition on disk

- 100Mb
- 100 000 – 1 000 000 cells

Reasons ? Make maintenance operations easier

- compaction
- repair
- bootstrap ...

Sub-partitioning techniques

PRIMARY KEY((sensor_id, day), dt))

→ 100 x 3600 x 24 = 8 640 000 cells on disks ✓

Sub-partitioning techniques

PRIMARY KEY((sensor_id, day), dt))

→ 100 x 3600 x 24 = 8 640 000 cells on disks ✓ □

But impact on queries:

- need to provide sensor_id & day for any query
- how to fetch data across N days ?

Data deletion and tombstones

```
DELETE FROM sensor_data  
  WHERE sensor_id = .. AND dt = ...
```

Logical deletion of data but:

- new physical "tombstone" column on disk
- disk space usage will increase !

The "tombstone" columns will be purged later by compaction process ...

Lab 3 : Hands-on Primary Keys

Vielen Dank!