



DataStax Hands-On Modelling

Negib Marhoul, Solution Engineer, DataStax

8. June 2017

Agenda

1	Storing data in Cassandra
2	Data Modelling, CQL basics
3	Hands-On Primary Keys

Storing data in Cassandra

Write request

custid : 4
name : Vincent

user_id : 4
name : Vincent

MEMTABLE

FLUSH

SSTABLE

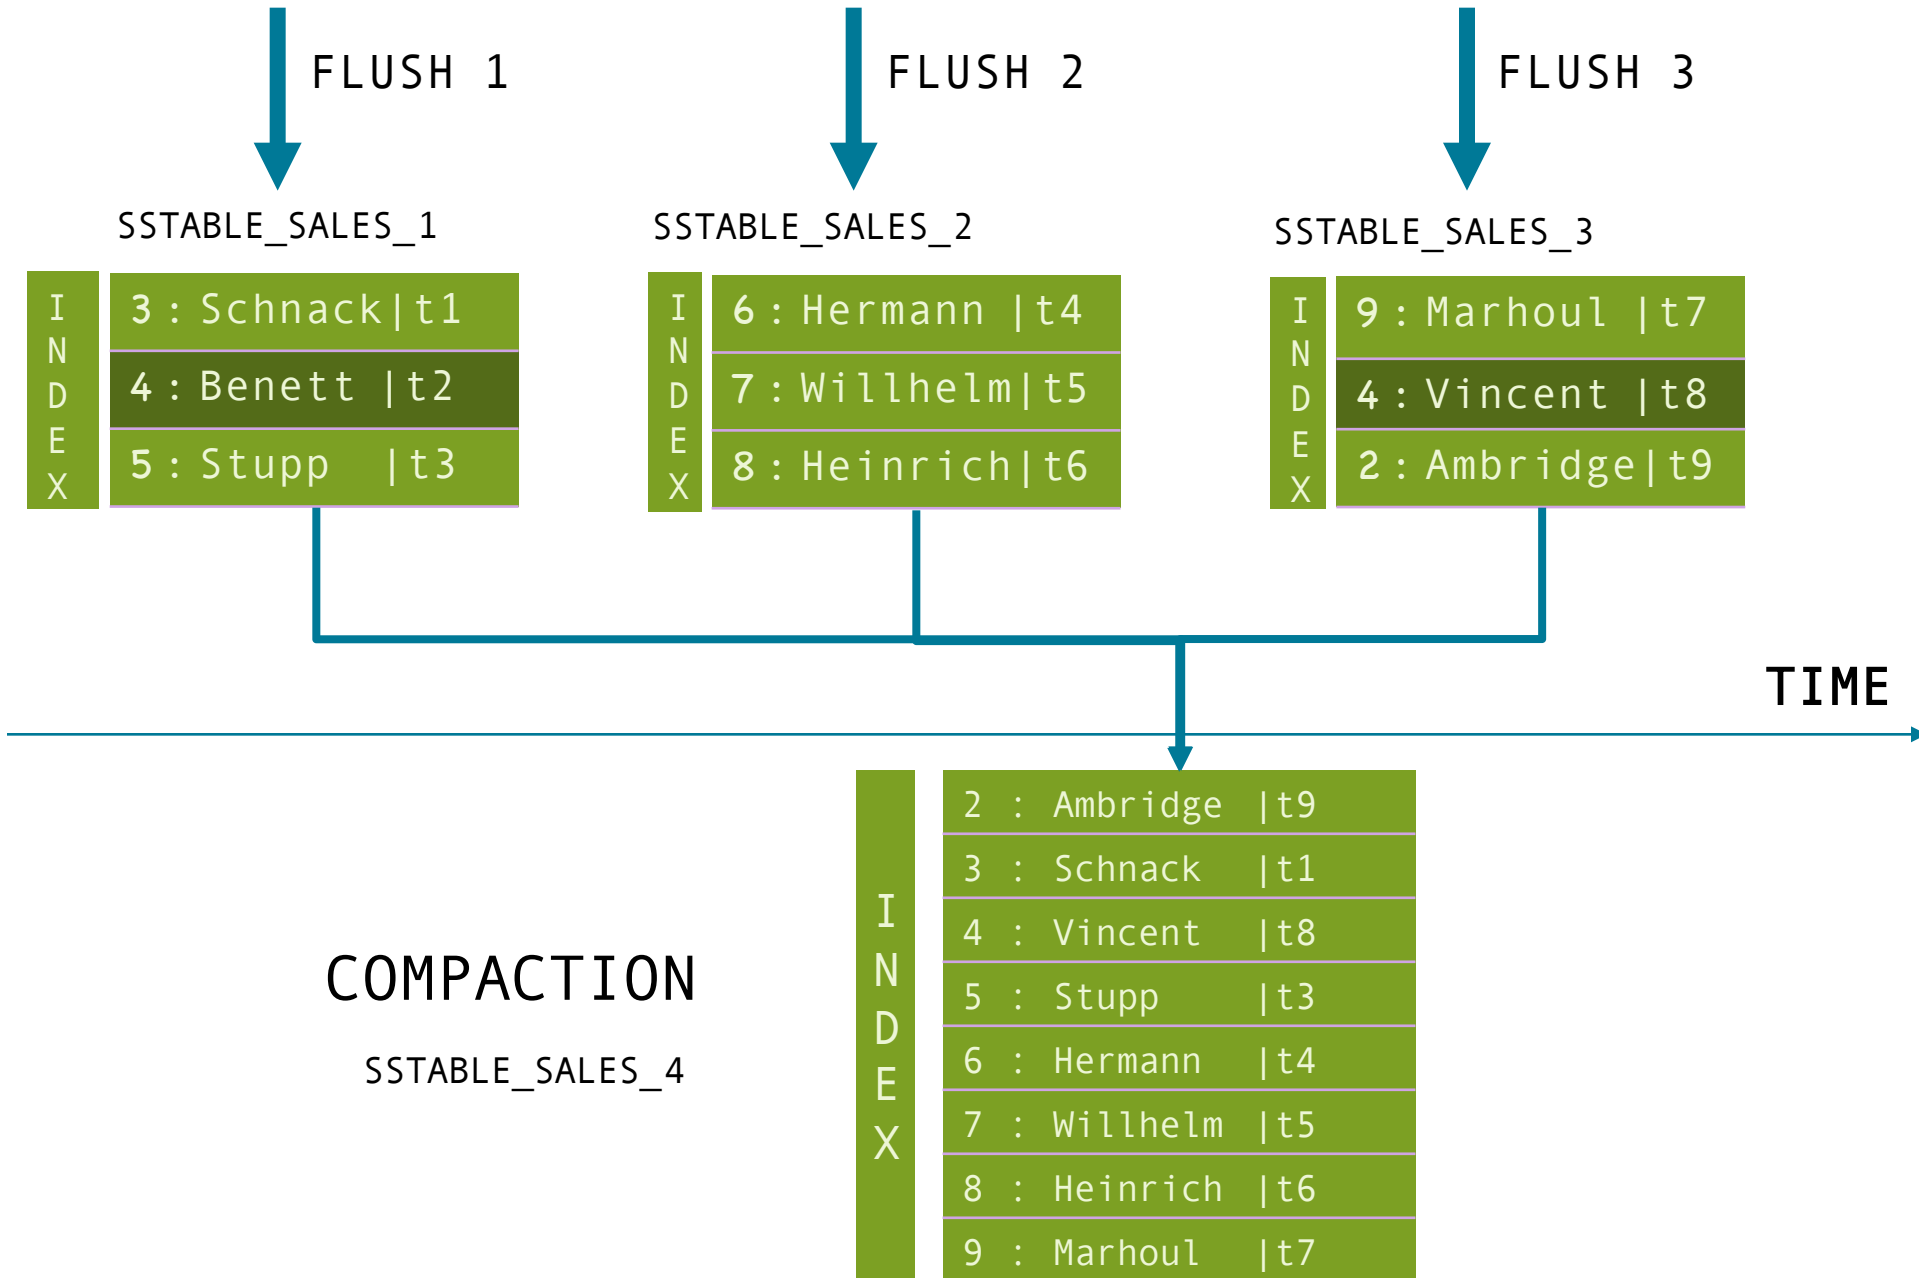
I
N
D
E
X

3	: Schnack
4	: Vincent
6	: Stupp
8	: Ambridge

IMMUTABLE
ORDERD

COMMIT LOG

user_id : 4
name : Vincent



Data Modeling Objectives

Data modeling objectives

1. Get your data **out of Cassandra**
2. Reduce query **latency**, make your queries **faster**
3. Avoid **disaster** in production

Data modeling methodology

Design by query

- first, know your **functional queries**
- then design the **table(s)** for direct access
- just **denormalize** if necessary
- Spread data evenly around the cluster
- Minimize the number of partitions read
- Make sure to take your read/update ratio into account when designing your schema

Output of design phase = **schema.cql**

Then start coding

Know your functional queries

Query:

`find users by id group by region and orderd by join date`

- Grouping by an attribute
- Ordering by an attribute
- Filtering based on some set of conditions
- Enforcing uniqueness in the result set

The partition key

Role

Partition key

- main entry point for query (INSERT/SELECT ...)
- help distribute/locate data on the cluster

No partition key = full cluster scan

How to choose correct partition key ?

Good partition column

- choose **functional** identifier
- high cardinality (lots of **distinct values**)

Query:

Find sales by session?


```
CREATE TABLE sales_by_customer (  
    custid int,  
    salesdt date,  
    revenue double,  
    discount double,  
    comment txt,  
    PRIMARY KEY(custid));
```

partition key (#partition)




Example of good partition key

```
CREATE TABLE sales_by_order(  
  custid int,  
  ... ,  
  PRIMARY KEY(custid))
```



Example of bad partition key

```
CREATE TABLE sales_by_cat(  
  category text,  
  ...,  
  PRIMARY KEY(category))
```



CRUD operations

```
INSERT INTO sales_by_customer (custid, salesdt, revenue) VALUES('1', '20160103', 799);
```

```
UPDATE sales_by_customer SET discount = 10 WHERE custid = 1;
```

```
DELETE sales_by_customer FROM sales_by_order WHERE custid = 1;
```

```
SELECT revenue FROM sales_by_customer WHERE custid = 1;
```

No Clustering Columns

PRIMARY KEY (custid)

// no duplicate primary keys

custid	salesdt	revenue	discount	comment
1	20160101	300	10	PlayStation4
2	20160102	699	25	iPhone 7
2	20160103	750	10	iPhone 7

```
SELECT * FROM sales_by_customer WHERE custid = '2';
```


Composite partition key

Multiple columns for partition key

- always **known in advance** (INSERT/SELECT ...)
- are **hashed together** to the same token value

```
CREATE TABLE sales_by_customer (  
    custid int,  
    salesdt date,  
    revenue double,  
    discount double,  
    comment txt,  
    PRIMARY KEY(custid, salesdt));
```

```
SELECT * FROM sales_by_customer WHERE custid = ... AND salesdt = ...
```

```
SELECT * FROM sales_by_customer WHERE custid = ... AND salesdt IN (xxx, yyy ...)
```

Compound Partition Key

PRIMARY KEY ((custid, salesdt)) // hash(custid, salesdt) → token

custid	salesdt	revenue	discount	comment
1	20160101	300	10	PlayStation4
2	20160102	699	25	IPhone 7
2	20160103	750	10	IPhone 7

```
SELECT * FROM sales_by_customer WHERE custid = 2 AND salesdt = '1/2/2016';
```

The clustering column(s)

Role

Clustering column(s)

- simulate 1 – N relationship
- and sort data (logically & on disk)

Clustered table (1 – N)

```
CREATE TABLE sales_by_customer (  
  custid  int,  
  salesdt date,  
  revenue int,  
  discount int,  
  comment text,  
  PRIMARY KEY(custid, salesdt));
```

partition key

clustering column
(sorted)

unicity

Recommended syntax

```
PRIMARY KEY((custid), salesdt);
```

Clustering Columns Create Wide Rows

PRIMARY KEY ((custid),salesdt)


// default sort and range queries

custid	salesdt		
	20160102		
1	719, 10, new customer iPhone7		
2	20160101	20160102	20160103
	719, 10, new customer iPhone7	45, 50, marketing campaign xyz	187, 10, sonyplaystation 3

```
SELECT * FROM sales_by_customer WHERE custid = 1 and salesdt > '1/1/2016';
```

What's Stored With Each Column?

custid	salesdt		
1	20160102		
	719, 10, new customer iPhone7		
2	20160101	20160102	20160103
	719, 10, new customer iPhone7	45, 50, marketing campaign xyz	187, 10, sonyplaystation 3



column name : "comment"
column value : "new customer iPhone7"
timestamp : 1353890782373000
TTL : 3600

Columns relationship and ordering

```
CREATE TABLE sales_by_customer (  
  custid    int,  
  salesdt   date,  
  revenue   int,  
  discount  int,  
  comment   text,  
  PRIMARY KEY((custid), salesdt))  
WITH CLUSTERING ORDER BY (dt ASC)
```

Diagram illustrating column relationships and ordering:

- `custid(1) <-----> (N) salesdt` (indicated by an arrow from the `custid` and `salesdt` columns in the table definition)
- `salesdt (1) <-----> (1) (revenue, discount, comment)` (indicated by an arrow from the `revenue`, `discount`, and `comment` columns in the table definition)

custid	salesdt		
1	20160102		
	719, 10, new customer iPhone7		
2	20160101	20160102	20160103
	719, 10, new customer iPhone7	45, 50, marketing campaign xyz	187, 10, sonyplaystation 3

```
SELECT * FROM sales_by_customer WHERE custid = 1 and salesdt > '1/1/2016';
```


Multiple clustering columns

```
CREATE TABLE sales_by_cust (  
  custid    int,  
  salesch   text,  
  salesdt   date,  
  revenue   int,  
  discount  int,  
  comment   text,  
  PRIMARY KEY((custid), salesch, salesdt))  
WITH CLUSTERING ORDER BY (salesch ASC, salesdt DESC)
```

→ salesch(1) <-----> (N) date

custid		salesch, salesdt	
1	online		
	20160102		
	719, 10, new customer iPhone7		
2	online	online	store
	20160101	20160102	20160103
	719, 10, new customer iPhone7	45, 50, marketing campaign xyz	187, 10, sonyplaystation 3

```
SELECT * FROM sales_by_customer  
WHERE custid = 1 AND  
salesch = "online" AND salesdt >= 1/2/16 AND salesdt <= 1/3/16;
```

Primary key summary

PRIMARY KEY((custid), salesch, salesdt)



Unicity of (custid, salesch, salesdt)

Primary key summary

PRIMARY KEY((custid), salesch, salesdt)



Used to locate node in the cluster

Used to locate partition in the node

Primary key summary

PRIMARY KEY((custid), salesch, salesdt)



Used to lookup rows in a partition

Used for data sorting and range queries

Lab 3 : Hands-on Primary Keys

Other critical details

Huge partitions

```
PRIMARY KEY((sensor_id), dt))
```

Data for the same sensor stay in the same partition on disk

Huge partitions

PRIMARY KEY((sensor_id), dt))

Data for the same sensor stay in the same partition on disk

If insert rate = 100/sec, how big is my partition after 1 year ?

→ $100 \times 3600 \times 24 \times 365 = 3\,153\,600\,000$ cells on disks

Huge partitions

PRIMARY KEY((sensor_id), dt))

Theoretical limit of # cells for a partition = 2×10^9

Practical limit for a partition on disk

- 100Mb
- 100 000 – 1 000 000 cells

Reasons ? Make maintenance operations easier

- compaction
- repair
- bootstrap ...

Sub-partitioning techniques

PRIMARY KEY((sensor_id, day), dt))

→ 100 x 3600 x 24 = 8 640 000 cells on disks ✓

Sub-partitioning techniques

PRIMARY KEY((sensor_id, day), dt))

→ 100 x 3600 x 24 = 8 640 000 cells on disks ✓ □

But impact on queries:

- need to provide sensor_id & day for any query
- how to fetch data across N days ?

Data deletion and tombstones

```
DELETE FROM sensor_data  
  WHERE sensor_id = .. AND dt = ...
```

Logical deletion of data but:

- new physical "tombstone" column on disk
- disk space usage will increase !

The "tombstone" columns will be purged later by compaction process ...

Lab 3 : Hands-on Primary Keys

Vielen Dank!